

New Statistics

Statistics with R
Basel R Bootcamp



April 2019

New Statistics

New Statistics = Old Statistics

Better studies

- Informative designs
- Sample size planning
- Registrations
- No -hacking
- Full reporting
- Replication

Better statistics

- Reporting uncertainty
- Lowering
- Bayesian statistics



from [amazon.com](https://www.amazon.com)


Replication crises

There are various **replication crises rumbling** in diverse academic fields.

Large-scale replication attempt in **Psychology** found that **only 36% were replicable**.

The low replicability is the result of **Questionable Research Practices**.

Similar assessments in **Medicine, Economics, Marketing, Social sciences**.



Hallmarks of good studies

- 1 Informative designs
- 2 **Sample size planning**
- 3 Registrations
- 4 **No p-hacking**
- 5 Full reporting
- 6 Replication

...and this is where we put the non-significant results.



adapted from [Someecards.com](https://www.someecards.com)


Sample size planning

Sample size should be planned such that there is **sufficient power to detect an effect**, if it is present.

Increasing sample size means **narrower sampling distributions** and **smaller decision errors**.

	Effect present	Effect absent
Significant result	$1-\beta$	α
Non-significant result	β	$1-\alpha$

Small N




Sample size planning

Sample size should be planned such that there is **sufficient power to detect an effect**, if it is present.

Increasing sample size means **narrower sampling distributions** and **smaller decision errors**.

	Effect present	Effect absent
Significant result	$1-\beta$	α
Non-significant result	β	$1-\alpha$



Sample size planning in R

Sample size should be planned such that there is **sufficient power to detect an effect**, if it is present.

Increasing sample size means **narrower sampling distributions** and **smaller decision errors**.

	Effect present	Effect absent
Significant result	$1-\beta$	α
Non- significant result	β	$1-\alpha$

```
# Load pwr package for pwr.t.test
library(pwr)

# N for large effect in t-test
pwr.t.test(sig.level = .05,
           power = .95,
           d = .8) # large effect
```

```
##
## Two-sample t test power calculation
##
##          n = 41.59
##          d = 0.8
##    sig.level = 0.05
##          power = 0.95
## alternative = two.sided
##
## NOTE: n is number in *each* group
```

Sample size planning in R

Sample size should be planned such that there is **sufficient power to detect an effect**, if it is present.

Increasing sample size means **narrower sampling distributions** and **smaller decision errors**.

	Effect present	Effect absent
Significant result	$1-\beta$	α
Non- significant result	β	$1-\alpha$

```
# Load pwr package for pwr.t.test
library(pwr)

# N for small effect in t-test
pwr.t.test(sig.level = .05,
           power = .95,
           d = .2) # small effect
```

```
##
##      Two-sample t test power calculation
##
##              n = 650.7
##              d = 0.2
##      sig.level = 0.05
##              power = 0.95
##      alternative = two.sided
##
## NOTE: n is number in *each* group
```


p-hacking

Ronald Coase

from fivethirtyeight.com

Hack Your Way To Scientific Glory

You're a social scientist with a hunch: The U.S. economy is affected by whether Republicans or Democrats are in office. Try to show that a connection exists, using real data going back to 1948. For your results to be publishable in an academic journal, you'll need to prove that they are "statistically significant" by achieving a low enough p-value.

1 CHOOSE A POLITICAL PARTY

Republicans

Rep.

Democrats

Dem.

2 DEFINE TERMS

Which politicians do you want to include?

Politicians

☐ **Presidents** Pres.

☒ **Governors** Govs.

☒ **Senators** Sens.

☐ **Representatives** Reps.

How do you want to measure economic performance?

Economic factors

☐ **Employment** Jobs

9 / 20

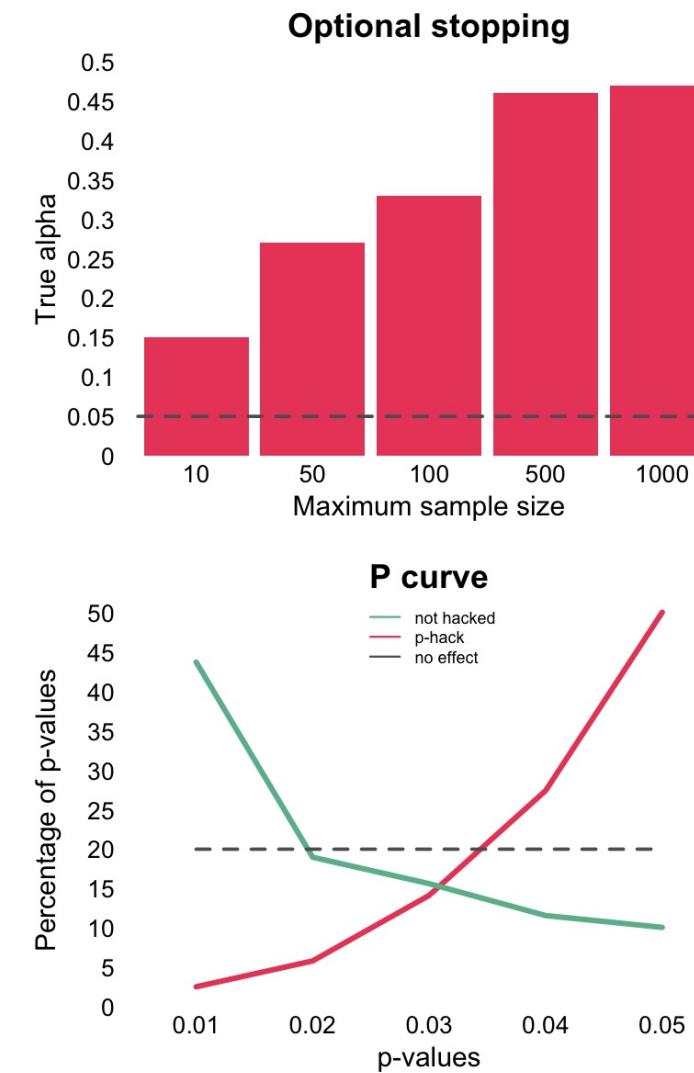
p-hacking

p-hacking is the misuse of data analysis to find patterns in data that can be presented as statistically significant when in fact there is no real underlying effect.

Wikipedia

Optional Stopping - Stopping data collection once significance is reached.

HARKing - Hypothesizing After the Data are Known. Committed when presenting any non-planned analyses, including the introduction of covariates.



Dos and Don'ts



Columbus looking for India, from [history.com](https://www.history.com)

Everything.

Present your result as confirmatory.



Villemard vision for 2000, from [sadanduseless.com](https://www.sadanduseless.com)

Make predictions.
Predetermine sample size.
Predetermine analysis plan.
Register.

Engage in non-planned analyses.

New Statistics

Problems with p -values

They are **difficult to interpret** as measures of evidence.


They are used for **arbitrary binarization** via comparison to α .

New statistics

Focus on **estimation**:
rather than

Communication of **uncertainty**:
rather than

$$\text{Confidence Interval (CI)} = \bar{x} \pm t_{1-\alpha} \sigma_{\bar{x}}$$



Confidence interval


Confidence intervals essentially are **rearranged significance tests**.

t-test $\frac{|\bar{x}|}{\sigma_{\bar{x}}} > t_{1-\alpha}$

Step 1 $|\bar{x}| > t_{1-\alpha} \sigma_{\bar{x}}$

Step2 $|\bar{x}| - t_{1-\alpha} \sigma_{\bar{x}} > 0$

CI $\bar{x} - t_{1-\alpha} \sigma_{\bar{x}} > 0$
 $\bar{x} + t_{1-\alpha} \sigma_{\bar{x}} < 0$



Confidence interval

Confidence intervals essentially are **rearranged significance tests**.

$$\text{t-test} \quad \frac{|\bar{x}|}{\sigma_{\bar{x}}} > t_{1-\alpha}$$

$$\text{Step 1} \quad |\bar{x}| > t_{1-\alpha} \sigma_{\bar{x}}$$

$$\text{Step2} \quad |\bar{x}| - t_{1-\alpha} \sigma_{\bar{x}} > 0$$

$$\text{CI} \quad \begin{aligned} \bar{x} - t_{1-\alpha} \sigma_{\bar{x}} &> 0 \\ \bar{x} + t_{1-\alpha} \sigma_{\bar{x}} &< 0 \end{aligned}$$

```
# standard confidence interval
N      <- length(t_1)
Delta  <- mean(t_1 - t_2)
SE      <- sd(t_1 - t_2) / sqrt(N)
Delta + SE * qt(.95, N - 1) * c(-1, 1)
```

```
## [1] -1.1480 -0.7421
```

Confidence interval

Confidence intervals essentially are **rearranged significance tests**.

$$\text{t-test} \quad \frac{|\bar{x}|}{\sigma_{\bar{x}}} > t_{1-\alpha}$$

$$\text{Step 1} \quad |\bar{x}| > t_{1-\alpha} \sigma_{\bar{x}}$$

$$\text{Step2} \quad |\bar{x}| - t_{1-\alpha} \sigma_{\bar{x}} > 0$$

$$\text{CI} \quad \begin{aligned} \bar{x} - t_{1-\alpha} \sigma_{\bar{x}} &> 0 \\ \bar{x} + t_{1-\alpha} \sigma_{\bar{x}} &< 0 \end{aligned}$$

```
# bootstrapped confidence interval
bf <- function(x,ind) {
  sum(x[ind])/length(x[ind])
}
boot_res <- boot(t_1 - t_2, bf, 1000)
boot.ci(boot_res)
```

```
## Warning in boot.ci(boot_res): bootstrap
## variances needed for studentized intervals
```

```
## BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS
## Based on 1000 bootstrap replicates
##
```

```
## CALL :
## boot.ci(boot.out = boot_res)
##
```


```
## Intervals :
## Level      Normal      Basic
## 95%   (-1.1777, -0.7113 )  (-1.1896, -0.7136 )
```

Bayesian statistics

Bayesian statistics extends classic (Frequentist) statistics by a **prior** distribution, specifying the prior **probability of the hypotheses** before seeing the data.

The prior permits calculation of a **true -value**, the **posterior probability**.

Posterior	Likelihood	Prior
Probability of the hypothesis given the data: <i>A real p-value</i>	Or the classic <i>p-value</i>	Initial belief in the hypothesis
↓	↓	↓
$p(H D) = \frac{p(D H)p(H)}{p(D)}$		
		↗
Normalization Sum of numerators $p(D) = \sum p(D H_i)p(H_i)$		



from xkcd.com


Bayesian statistics

For two reasons Bayesian statistics has long been rejected...

- 1 **Computational needs** to calculate $p(D)$
- 2 **Subjectivity** in choosing the prior

However,...

- 1 Computers have become **fast enough**
- 2 Statistical inference is **necessarily subjective**
- 3 The influence of priors can be limited through **ignorance priors**



from xkcd.com

Bayesian statistics

The `rstanarm` and `BayesFactor` packages make Bayesian stats very easy to use.

Function	Package	Description
<code>stan_glm</code> , <code>stan_glmer</code>	<code>rstanarm</code>	Bayesian (mixed) regression
<code>ttestBF</code> , <code>anovaBF</code>	<code>BayesFactor</code>	Standard h-tests
<code>lmBF</code>	<code>BayesFactor</code>	Bayesian (mixed) regressions

```
# Bayesian stats with rstanarm
library(rstanarm)

stan_glm(formula = income ~ height,
          data = baselers)
```

```
## stan_glm
## family:      gaussian [identity]
## formula:     income ~ height
## observations: 300
## predictors:  2
## -----
##              Median MAD_SD
## (Intercept) 9292.9 2188.2
## height      -10.9  12.8
##
## Auxiliary parameter(s):
##              Median MAD_SD
## sigma 2635.5  107.2
##
## Sample avg. posterior predictive distribution of y:
##              Median MAD_SD
## mean_PPD 7439.5  211.7
##
## -----
## * For help interpreting the printed output see ?print.stanreg
## * For info on the priors used see ?prior_summary.stanreg
```

Bayesian statistics

The `rstanarm` and `BayesFactor` packages make Bayesian stats very easy to use.

Function	Package	Description
<code>stan_glm</code> , <code>stan_glmer</code>	<code>rstanarm</code>	Bayesian (mixed) regression
<code>ttestBF</code> , <code>anovaBF</code>	<code>BayesFactor</code>	Standard h-tests
<code>lmBF</code>	<code>BayesFactor</code>	Bayesian (mixed) regressions

```
# Bayesian stats with BayesFactor
library(BayesFactor)

lmBF(formula = income ~ height,
      data = baselers)
```

```
## Bayes factor analysis
## -----
## [1] height : 0.1813 ±0%
##
## Against denominator:
##   Intercept only
## ---
## Bayes factor type: BFlinearModel, JZS
```

Practical